How BARCS Finds Important Genes

Chapter 2: The calculations

How BARCS fits the model and decides if an effect is real

Prepared for Hyun-Hwan Jeong and Kyu-Won Lee

BARCS version 0.2.1 · Reading level: grades 7 to 9

About this guide. This guide explains how the BARCS computer program works. You need only arithmetic, fractions, percents, powers and square roots. The guide teaches logarithms when it uses them. Both chapters use the same small set of numbers, so you can follow every step. The guide uses the writing rules of ASD-STE100 Simplified Technical English.

Contents

What you need from Chapter 1 **2**

Step 6: How BARCS finds the best β values **2**

The loop 2

The formulas 3

Iteration 1 in full 4

The next iterations 4

Step 7: How BARCS measures the biological noise **5**

Pearson residuals 5

The rule for ρ 5

Step 8: Is the effect real? **6**

The standard error 6

The t value 7

Degrees of freedom 7

The p-value 7

Step 9: Borrow information from other guides **8**

What moderation does to the test 9

Step 10: Check the results with control guides **9**

Step 11: From guides to genes **10**

When guides share noise 11

Step 12: Many genes and false discoveries **11**

Two special cases **12**

A guide with zero reads 12

A very large change 12

Summary: the full path **12**

Quiz for Chapter 2 **14**

Answers for Chapter 2 **17**

What you need from Chapter 1

This chapter continues the example of Chapter 1. Here is a short summary.

The data. One guide, four libraries:

Library	Group	Guide count K	Total count N
1	control	180	1,000,000
2	control	240	1,200,000
3	drug	120	1,000,000
4	drug	80	800,000

- **Proportion** $p = K/N$. Combined: 191 per million (control) and 111 per million (drug). The fold change is 0.58.
- **Log-odds** $= \ln\Big(p/(1-p)\Big)$: -8.56 (control) and -9.10 (drug).
- **Noise.** The variance of a count is $N \times \mu \times (1-\mu) \times \Big(1 + (N-1) \times \rho\Big)$. Here μ is the expected proportion and ρ measures the biological noise.
- **Design matrix** X: one row for each library, one column for each β. Our X has the rows $(1,0)$, $(1,0)$, $(1,1)$, $(1,1)$.
- **Prediction.** Multiply a row of X by the β values and add. The result is the predicted log-odds η. Then $\mu = e^{\eta}/(1+e^{\eta})$.
- β_0 is the control log-odds. β_1 is the effect of the drug.

Step 6: How BARCS finds the best β values

BARCS cannot find β_0 and β_1 in one calculation. It uses a loop of "guess, check, improve". The name of this method is **iteratively reweighted least squares** (IRLS). Each pass through the loop is an **iteration**.

The loop

1. **Guess.** Start with the simple log-odds of each group (Step 3 in Chapter 1).
2. **Predict.** Use the design matrix to predict η and μ for each library (Step 5 in Chapter 1).
3. **Measure the noise.** Find ρ from the differences between the real and the expected counts (Step 7).
4. **Weigh.** Give each library a **weight**. A library with less noise gets a larger weight.

5. **Correct.** Calculate a corrected value, the **working value**, for each library.
6. **Improve.** Take a weighted average of the working values in each group. These averages give the new β values.
7. **Repeat** from part 2 until the β values stop changing.

The formulas

The weight of library i is

$$w_i = \frac{N_i \times \mu_i \times (1 - \mu_i)}{1 + (N_i - 1) \times \rho}.$$

The top is the counting noise turned upside down. The bottom is the extra biological noise from Step 4 in Chapter 1. So more noise gives a smaller weight.

The working value of library i is

$$\text{working value}_i = \eta_i + \frac{K_i/N_i - \mu_i}{\mu_i \times (1 - \mu_i)}.$$

It is the current prediction η_i plus a correction. The correction is large when the real proportion K_i/N_i is far from the expected proportion μ_i.

For our design matrix, the new β values are weighted averages:

$$\beta_0 = \frac{w_1\, v_1 + w_2\, v_2}{w_1 + w_2}, \qquad \beta_0 + \beta_1 = \frac{w_3\, v_3 + w_4\, v_4}{w_3 + w_4},$$

where v_i is the working value of library i. (For other design matrices, BARCS solves a small set of equations instead. The idea is the same.)

Iteration 1 in full

Worked example 10: iteration 1

Guess. From Step 3 in Chapter 1, the control log-odds is -8.563522 and the drug log-odds is -9.104869. So $\beta_0 = -8.563522$ and $\beta_1 = -9.104869 - (-8.563522) = -0.541347$.

Predict. The control libraries get $\mu = 0.000190909$. The drug libraries get $\mu = 0.000111111$.

Measure the noise. Step 7 gives $\rho = 0.0000003841$ for these predictions.

Weigh library 1. Top: $1{,}000{,}000 \times 0.000190909 \times 0.999809 = 190.87$.
Bottom: $1 + 999{,}999 \times 0.0000003841 = 1.3841$.
Weight: $w_1 = 190.87 \div 1.3841 = 137.90$.

Correct library 1. $K_1/N_1 - \mu_1 = 0.000180 - 0.000190909 = -0.0000109$.
$\mu_1(1 - \mu_1) = 0.000190873$.
Correction: $-0.0000109 \div 0.000190873 = -0.0572$.
Working value: $v_1 = -8.5635 + (-0.0572) = -8.6207$.

The same calculations for all four libraries give:

Library	Weight w_i	Working value v_i	$w_i \times v_i$
1 (control)	137.905	−8.62068	−1188.84
2 (control)	156.785	−8.51589	−1335.16
3 (drug)	80.269	−9.02486	−724.42
4 (drug)	67.988	−9.20488	−625.82

Improve.
Control: $\beta_0 = (-1188.84 - 1335.16) \div (137.905 + 156.785) = -2524.00 \div 294.69 = -8.564929$.
Drug: $\beta_0 + \beta_1 = (-724.42 - 625.82) \div (80.269 + 67.988) = -1350.24 \div 148.257 = -9.107414$.
So $\beta_1 = -9.107414 - (-8.564929) = -0.542486$.

The next iterations

BARCS repeats the loop with the new values. The changes become very small:

After iteration	β_0	β_1
0 (the guess)	−8.563522	−0.541347
1	−8.564929	−0.542486
2	−8.564934	−0.542492
3	−8.564934	−0.542492

After iteration 3, the values do not change in the first six decimal places. BARCS stops. The final values are $\beta_0 = -8.565$ and $\beta_1 = -0.542$.

Remember

The weighted average gives a slightly different answer from the simple log-odds of Step 3 in Chapter 1 (-0.5425 instead of -0.5413). The difference comes from the weights. Library 2 has the most reads and the least relative noise. So it counts the most.

Step 7: How BARCS measures the biological noise

Step 4 in Chapter 1 used $\rho = 0.000000386$. This section shows where the number comes from.

Pearson residuals

For each library, BARCS subtracts the expected count from the real count. It divides the result by the standard deviation from Step 4 in Chapter 1. The result is a **Pearson residual**:

$$\text{Pearson residual}_i = \frac{K_i - N_i \times \mu_i}{\sqrt{N_i \times \mu_i \times (1 - \mu_i) \times \Big(1 + (N_i - 1) \times \rho\Big)}}.$$

A Pearson residual tells how many standard deviations a library is from its prediction.

The rule for ρ

If the noise model is correct, the squares of the Pearson residuals add up to about the degrees of freedom. So BARCS chooses ρ to make this statement exactly true:

$$\text{sum of the squared Pearson residuals} = \text{df}.$$

Worked example 11: start with $\rho = 0$

Use the expected counts from Step 5 in Chapter 1. With $\rho = 0$, the bottom of each residual is the counting noise only.

Library	$(K - N \times \mu)^2$	$N \times \mu \times (1 - \mu)$	Squared residual
1	$(-10.64)^2 = 113.21$	190.60	0.594
2	$(+11.23)^2 = 126.16$	228.72	0.552
3	$(+9.17)^2 = 84.14$	110.82	0.759
4	$(-8.66)^2 = 75.03$	88.65	0.846
		Sum	2.751

The sum is 2.751. The degrees of freedom are 2. The sum is too big. So there is more noise than counting noise alone. We need $\rho > 0$.

Worked example 12: guess and check for ρ

A larger ρ makes the bottom of each residual larger. So the sum goes down. BARCS tries values until the sum equals 2:

ρ	Sum of squared residuals
0	2.751
0.0000001	2.506
0.0000002	2.302
0.0000003	2.129
0.000000386	2.000
0.0000005	1.851

The answer is $\rho = 0.000000386$. If the sum is already 2 or less at $\rho = 0$, BARCS uses $\rho = 0$.

The ratio $2.751 \div 2 = 1.376$ is also useful. It is the **variance inflation**. It tells that the guide has about 1.38 times more noise than counting noise alone. Step 9 uses this number.

Step 8: Is the effect real?

The effect is $\beta_1 = -0.542$. But replicates are noisy. Maybe the true effect is 0 and the noise made the drop. BARCS tests this question in four parts.

The standard error

The **standard error** (SE) tells how much the estimate of β_1 can change because of noise. A small SE means a precise estimate.

For our design matrix, the standard error comes from the weights of Step 6. Add the weights of each group:

$$W_{\text{control}} = w_1 + w_2, \qquad W_{\text{drug}} = w_3 + w_4.$$

Then

$$\text{SE}(\beta_1) = \sqrt{\frac{1}{W_{\text{control}}} + \frac{1}{W_{\text{drug}}}}.$$

More weight means less noise, so the SE becomes smaller. (For other design matrices, BARCS uses the inverse of the matrix $X^\top W X$. For our design, that calculation gives exactly this formula.)

Worked example 13: the standard error

The final weights are 137.524, 156.322, 79.955 and 67.737.
$W_{\text{control}} = 137.524 + 156.322 = 293.846$.
$W_{\text{drug}} = 79.955 + 67.737 = 147.692$.
$1 \div 293.846 = 0.003403$.
$1 \div 147.692 = 0.006771$.
Sum: $0.003403 + 0.006771 = 0.010174$.
SE $= \sqrt{0.010174} = 0.101$.
The control group has about twice the weight of the drug group. So most of the uncertainty comes from the drug group.

The t value

The **t value** compares the effect with its standard error:

$$t = \frac{\beta_1}{\text{SE}}.$$

A t value far from 0 means that the effect is large compared with the noise.

Degrees of freedom

The **degrees of freedom** (df) tell how much information we have about the noise:

$$\text{df} = (\text{rows of } X) - (\text{columns of } X) = (\text{number of libraries}) - (\text{number of } \beta \text{ values}).$$

Fewer degrees of freedom mean less information. Then BARCS needs a bigger t value before it accepts an effect as real.

Degrees of freedom	1	2	14	very many
$\lvert t \rvert$ needed for $p < 0.05$	12.71	4.30	2.14	1.96

The p-value

The **p-value** answers this question: "If the true effect is 0, how often does noise alone give a t value this far from 0?" A small p-value means that noise alone seldom gives such a result. Scientists often use 0.05 as a limit.

Worked example 14: testing the effect

1. $t = -0.542 \div 0.101 = -5.38$.
2. X has 4 rows and 2 columns. So $\text{df} = 4 - 2 = 2$.
3. The table shows that 2 degrees of freedom need $|t| > 4.30$. Our $|t| = 5.38$. It is larger than 4.30. So $p < 0.05$.
4. A computer gives the exact p-value: $p = 0.033$.

The p-value is less than 0.05. So the drop is probably not only noise.

Remember

The degrees of freedom come from the number of libraries, not from the number of reads. A library with 100 million reads still counts as one row of X.

Step 9: Borrow information from other guides

With only 2 degrees of freedom, the noise estimate of one guide is not reliable. But a screen has thousands of guides. Most guides have a similar amount of noise. So BARCS combines two estimates:

- the variance inflation of this guide (Step 7), and
- the typical variance inflation of guides with similar counts.

This method is **moderation**. BARCS takes a weighted average. Each estimate gets a number of "votes". The guide's own estimate gets its degrees of freedom as votes. The typical estimate gets the **prior degrees of freedom**, d_0. BARCS calculates d_0 and the typical value from all the guides in the screen.

$$\text{moderated inflation} = \frac{\text{df} \times \text{own inflation} + d_0 \times \text{typical inflation}}{\text{df} + d_0}.$$

Worked example 15: moderation with simple numbers

A guide has 3 degrees of freedom. Its own inflation is 4.0. The typical inflation is 1.5, with $d_0 = 12$ votes.

$$\text{moderated inflation} = \frac{3 \times 4.0 + 12 \times 1.5}{3 + 12} = \frac{12 + 18}{15} = \frac{30}{15} = 2.0.$$

The guide's value moves from 4.0 toward 1.5. It stops at 2.0.

What moderation does to the test

BARCS uses the moderated inflation in three places:

1. It multiplies the standard error by $\sqrt{\text{moderated inflation} \div \text{own inflation}}$.
2. It recalculates t with the new standard error.
3. It adds the d_0 votes to the degrees of freedom. So the test uses $\text{df} + d_0$.

Worked example 16: moderation for our guide

Our guide has own inflation 1.376 (Step 7) and 2 degrees of freedom. Suppose that the screen gives typical inflation 1.5 with $d_0 = 12$.
Moderated inflation: $(2 \times 1.376 + 12 \times 1.5) \div (2 + 12) = (2.751 + 18) \div 14 = 1.482$.
Factor for the SE: $\sqrt{1.482 \div 1.376} = \sqrt{1.077} = 1.038$.
New SE: $0.101 \times 1.038 = 0.105$.
New t: $-0.542 \div 0.105 = -5.18$.
New degrees of freedom: $2 + 12 = 14$.
With 14 degrees of freedom, $|t|$ must be larger than 2.14 for $p < 0.05$. Our $|t| = 5.18$.
The new p-value is 0.00014.
The SE became slightly larger, because the typical guide is a little noisier than our guide. But the degrees of freedom grew from 2 to 14. So the test became much more powerful.

BARCS uses moderation when it has 50 or more usable guides.

Step 10: Check the results with control guides

A good screen contains **non-targeting control guides**. These guides switch off no gene. So their true effect is 0.

If the noise model is correct, about 5% of control guides get $p < 0.05$ by chance. If many more control guides get $p < 0.05$, the noise model is too small. Then BARCS makes all the tests wider.

Worked example 17: calibration with control guides

A screen has 100 control guides. We expect about 5 of them to get $p < 0.05$. But 15 of them get $p < 0.05$. The noise is too small.

BARCS finds a **scale factor** that makes the control guides behave correctly. It sorts the 100 control guides by the size of $|t|$ (the value without its sign). It takes the value that 95 of the 100 guides stay below. Suppose that this value is 2.47. With many degrees of freedom, the t table says that this value should be 1.96. So the scale factor is $2.47 \div 1.96 = 1.26$. BARCS divides every t value by 1.26.

A gene with $t = 3.5$ now gets $t = 3.5 \div 1.26 = 2.78$. Its p-value becomes larger. This change protects you from false results.

The scale factor is never less than 1. So calibration can only make the results more careful.

Step 11: From guides to genes

Each gene usually has several guides, often 3 to 6. Each guide gives its own result. BARCS combines them into one result for the gene. It uses the **Stouffer method** in three parts.

1. Change each guide's p-value into a **z score**. A z score tells how many standard deviations a result is from 0. The sign of the z score follows the sign of the effect.
2. Add the z scores.
3. Divide the sum by the square root of the number of guides, m.

$$Z_{\text{gene}} = \frac{z_1 + z_2 + \cdots + z_m}{\sqrt{m}}.$$

To change a p-value into a z score, use the table of the normal distribution. A two-sided p-value of 0.05 belongs to $z = 1.96$. Some other pairs:

Two-sided p-value	0.317	0.134	0.046	0.012	0.0027
z score	1.0	1.5	2.0	2.5	3.0

If the effect is negative (the guide drops), the z score gets a minus sign.

Worked example 18: combining three guides

A gene has $m = 3$ guides. Their p-values are 0.046, 0.134 and 0.012. Their z scores are 2.0, 1.5 and 2.5.

$$Z_{\text{gene}} = \frac{2.0 + 1.5 + 2.5}{\sqrt{3}} = \frac{6.0}{1.732} = 3.46.$$

This Z gives $p = 0.00053$ for the gene. No single guide is that strong. Together, the three guides give strong evidence.

When guides share noise

Sometimes the guides of one gene share noise. Then they agree more than they should by chance. BARCS measures this sharing with the **guide correlation**, r. It makes the bottom of the fraction bigger:

$$Z_{\text{gene}} = \frac{z_1 + z_2 + \cdots + z_m}{\sqrt{m + m(m-1)\,r}}.$$

Worked example 19: shared noise

Use the same three guides and $r = 0.1$. The bottom becomes $\sqrt{3 + 3 \times 2 \times 0.1} = \sqrt{3.6} = 1.897$. So $Z_{\text{gene}} = 6.0 \div 1.897 = 3.16$, and $p = 0.0016$. The evidence is still strong, but less strong than before.

Step 12: Many genes and false discoveries

A screen tests about 20,000 genes. Suppose that no gene has an effect. If we use the limit $p < 0.05$, about 5% of the genes still pass by chance. That is about 1,000 false results.

So BARCS uses the **false discovery rate** (FDR). The FDR is the expected fraction of false results in your list of hits. BARCS calculates it with the **Benjamini–Hochberg method**:

1. Sort the genes from the smallest p-value to the largest.
2. Give each gene a rank: 1, 2, 3, and so on.
3. For each gene, calculate $p \times$ (number of genes) $\div$ rank.
4. Go from the bottom of the list to the top. If a value is larger than the value below it, replace it with the value below it.

Worked example 20: false discovery rate for five genes

Gene	Rank	p	$p \times 5 \div$ rank	FDR (after part 4)
A	1	0.001	$0.001 \times 5 \div 1 = 0.005$	0.005
B	2	0.008	$0.008 \times 5 \div 2 = 0.020$	0.020
C	3	0.039	$0.039 \times 5 \div 3 = 0.065$	0.051 (from D)
D	4	0.041	$0.041 \times 5 \div 4 = 0.051$	0.051
E	5	0.600	$0.600 \times 5 \div 5 = 0.600$	0.600

With the limit FDR < 0.05, genes A and B are hits. With the limit FDR < 0.10, genes A, B, C and D are hits.

Two special cases

A guide with zero reads

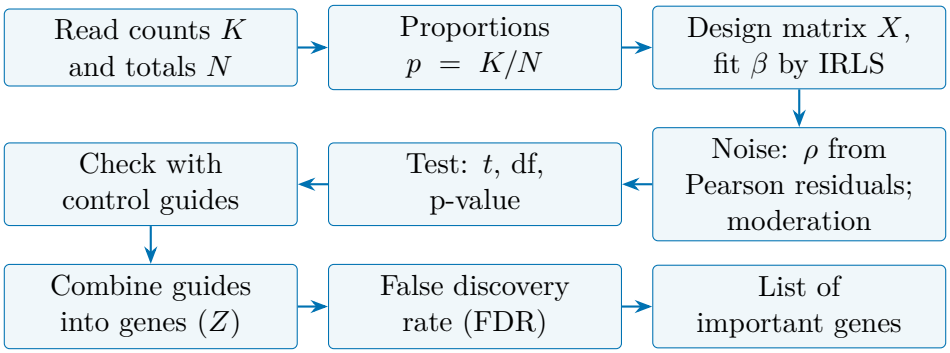
Sometimes a guide has 0 reads in all control libraries. Then the log-odds of the control group is $\ln(0)$. This value does not exist, because no power of e equals 0. Older versions of BARCS could not fit such a guide. So they removed it.

BARCS version 0.2.1 adds a very small number to the counts first. This number is less than one read. The method is the **Firth correction**. Now the guide gets a value, and BARCS keeps it.

A very large change

Sometimes a guide grows by hundreds of times. The usual test from Step 8 can then fail. The standard error grows faster than the effect, so the t value becomes small. BARCS version 0.2.1 has a second test, the **likelihood-ratio test**. It compares how well the model fits with the effect and without the effect. The option `test = "auto"` uses this test only for guides with very large changes.

Summary: the full path



All symbols in one table

Symbol	Meaning	Value in our example
K	read count of the guide in one library	180, 240, 120, 80
N	total read count of one library	1.0, 1.2, 1.0, 0.8 million
p	proportion, K/N	191 and 111 per million
fold change	$p_{\text{drug}} \div p_{\text{control}}$	0.58
$\log_2$	number of doublings	$\log_2(0.58) = -0.78$
ln	natural logarithm (powers of $e = 2.718$)	$\ln(0.58) = -0.54$
log-odds	$\ln(p/(1-p))$	-8.56 (control)
μ	expected proportion	0.000191 (control)
ρ	biological noise	0.000000386
X	design matrix: a row per library, a column per β	4 rows, 2 columns
β_0	control log-odds in the model	-8.565
β_1	effect of the drug	-0.542
η	predicted log-odds of a library (row of X times β)	-8.565, -9.107
w	weight of a library in the fit	137.5 (library 1)
working value	prediction plus correction (Example 10)	-8.621 (library 1)
Pearson residual	distance from the prediction in SDs	-0.65 (library 1)
inflation	noise compared with counting noise	1.376
SE	standard error of the effect	0.101
t	effect divided by SE	-5.38
df	libraries minus β values	$4 - 2 = 2$
p-value	chance of this result if the effect is 0	0.033
d_0	prior degrees of freedom (votes)	12 (Example 15)
z, Z	z score of a guide, of a gene	3.46 (Example 18)
m	number of guides for a gene	3
r	guide correlation (shared noise)	0.1 (Example 19)
FDR	expected fraction of false hits	0.005 for gene A

Note: the letter p has two uses in this guide. Alone, p is a proportion. In the word "p-value", it is a probability. The text always tells you which one it is.

Quiz for Chapter 2

Choose the best answer. The answers and calculations are on the next pages.

Q1. An experiment has 6 libraries. The model has 3 β values. How many degrees of freedom does it have?

(A) 3
(B) 2
(C) 6
(D) 9

Q2. An effect is $\beta_1 = -0.8$ with SE $= 0.2$. What is the t value?

(A) -0.16
(B) -1.0
(C) -0.6
(D) -4.0

Q3. A guide has df $= 2$ and its own noise 5.0. The typical noise is 2.0 with $d_0 = 8$. What is the moderated noise?

(A) 2.0
(B) 3.5
(C) 5.0
(D) 2.6

Q4. A screen has 200 control guides. About how many should get $p < 0.05$ by chance?

(A) 0
(B) 5
(C) 10
(D) 50

Q5. A gene has 4 guides, each with $z = 1.0$. What is Z_{gene} (no shared noise)?

(A) 1.0
(B) 2.0
(C) 4.0
(D) 0.5

Q6. Three genes have p-values 0.01, 0.02 and 0.30. What is the Benjamini–Hochberg value $p \times 3 \div$ rank for the second gene?

(A) 0.02

(B) 0.06

(C) 0.30

(D) 0.03

Q7. Why does BARCS 0.2.1 add a very small number to the counts of some guides?

(A) Because ln(0) does not exist when a group has zero reads.

(B) To make all guides look important.

(C) To increase the library totals.

(D) To remove control guides.

Q8. Which action gives BARCS more information about biological noise?

(A) Read each library more deeply.

(B) Use fewer guides.

(C) Add more replicate libraries.

(D) Use a larger p-value limit.

Q9. Two control libraries have weights 100 and 300. Their working values are -8.0 and -9.0. What is the weighted average?

(A) -8.5

(B) -8.25

(C) -8.75

(D) -9.0

Q10. The control weights add up to $W_{\text{control}} = 100$. The drug weights add up to $W_{\text{drug}} = 25$. What is $\text{SE}(\beta_1)$?

(A) 0.05

(B) 0.224

(C) 0.5

(D) 1.25

Q11. At $\rho = 0$, the squared Pearson residuals of a guide add up to 1.5. The model has 2 degrees of freedom. Which ρ does BARCS use?

(A) $\rho = 0$

(B) $\rho = 1.5$

(C) $\rho = 0.75$

(D) $\rho = 2$

Q12. Of 100 control guides, 95 have $|t|$ below 2.94. The t table says 1.96. What is the scale factor?

(A) 0.67

(B) 1.0
(C) 1.5
(D) 4.9

Answers for Chapter 2

Q1. **A.** $\text{df} = 6 - 3 = 3$ (Step 8).

Q2. **D.** $t = -0.8 \div 0.2 = -4.0$.

Q3. **D.** $(2 \times 5.0 + 8 \times 2.0) \div (2 + 8) = (10 + 16) \div 10 = 26 \div 10 = 2.6$ (Step 9).

Q4. **C.** 5% of $200 = 0.05 \times 200 = 10$ (Step 10).

Q5. **B.** $(1.0 + 1.0 + 1.0 + 1.0) \div \sqrt{4} = 4.0 \div 2 = 2.0$ (Step 11).

Q6. **D.** The second gene has rank 2. So $0.02 \times 3 \div 2 = 0.03$ (Step 12).

Q7. **A.** With zero reads, the log-odds is $\ln(0)$, which does not exist. The Firth correction fixes this problem.

Q8. **C.** Degrees of freedom come from libraries. Only more replicates add information about biological noise (Step 4 in Chapter 1 and Step 8).

Q9. **C.** $(100 \times (-8.0) + 300 \times (-9.0)) \div (100 + 300) = (-800 - 2700) \div 400 = -3500 \div 400 = -8.75$ (Step 6).

Q10. **B.** $\sqrt{1/100 + 1/25} = \sqrt{0.01 + 0.04} = \sqrt{0.05} = 0.224$ (Step 8).

Q11. **A.** The sum (1.5) is already 2 or less. So BARCS uses $\rho = 0$ (Step 7).

Q12. **C.** $2.94 \div 1.96 = 1.5$ (Step 10).